Alex Morgan

San Francisco, CA | alex.morgan@email.com | (555) 234-5678 | linkedin.com/in/alexmorgan-ai | github.com/alexmorgan

PROFESSIONAL SUMMARY

Results-driven Machine Learning Engineer with 3+ years of experience designing, training, and deploying production AI systems. Specializes in PyTorch, scikit-learn, LangChain, and RAG architectures. Experienced in building microservices with FastAPI and deploying containerized pipelines with Docker on AWS.

EDUCATION

Master of Science in Computer Science (Machine Learning Specialization)

Stanford University | 2020 – 2022

Bachelor of Science in Computer Science

University of California, Berkeley | 2016 – 2020

TECHNICAL SKILLS

• Programming Languages: Python, SQL, C++, Bash

• Machine Learning & AI: Machine Learning, Deep Learning, scikit-learn, PyTorch, TensorFlow, Pandas, NumPy, NLP, LLM, LangChain, LangGraph, RAG, Hugging Face, Transformers, Vector Databases, ChromaDB, FAISS

• Cloud & DevOps: Docker, AWS, Git, GitHub Actions, CI/CD, Linux

• Web & APIs: FastAPI, REST API, Flask

• Databases: PostgreSQL, Redis, MongoDB

PROFESSIONAL EXPERIENCE

Machine Learning Engineer | Apex AI Technologies | 2022 – Present

• Architected and deployed an enterprise Retrieval-Augmented Generation (RAG) pipeline using LangChain, ChromaDB, and OpenAI LLMs, cutting customer support response time by 42%.

• Built feature engineering pipelines and trained XGBoost and PyTorch neural network classifiers on tabular and text datasets exceeding 5M records with scikit-learn and Pandas.

• Packaged models as high-throughput REST APIs using FastAPI, containerized with Docker, and automated deployment via GitHub Actions on AWS ECS.

• Implemented TF-IDF and Cosine Similarity baseline search modules for semantic document retrieval prior to dense embedding vector ranking.

AI Research Intern | Cognition Labs | Summer 2021

• Researched few-shot prompt optimization for Large Language Models and fine-tuned Transformer models with PyTorch and Hugging Face.

• Developed data cleaning pipelines with NumPy, Pandas, and SciPy for multi-modal dataset analysis.

PROJECTS

• MedAssist RAG Agent: Multi-agent clinical document assistant using LangGraph, Python, FastAPI, and vector search over PubMed abstracts.

• Real-time Fraud Detection Engine: Stream processing ML service utilizing scikit-learn, Redis, and Docker for real-time anomaly detection.

CERTIFICATIONS

- AWS Certified Machine Learning - Specialty (2023)
- Deep Learning Specialization - DeepLearning.AI (Coursera)